

Sycophancy Impact Card													
Cheng et al. (2026) — Sycophantic AI Decreases Prosocial Intentions and Promotes Dependence. Science 391(6792), eaec8352.													
Behavior(s) measured	• Social validation: affirming users’ actions in interpersonal conflict scenarios, even when those actions involve deception, illegality, or harm (49% more than human baselines across 11 models) • Lack of pushback: failing to challenge one-sided or morally compromised self-accounts • Flattery / uncritical affirmation: validating users’ moral self-assessments without engaging critical perspective The paper introduces the term social sycophancy for this cluster of behaviors, distinct from the opinion-conformity behaviors dominant in prior work.												
Referent	Person primary Position secondary The primary target of sycophancy is the user themselves — their moral standing, character, and judgment — rather than their stated opinions on external facts. Position sycophancy is present secondarily, as the model also endorses the user’s framing of who is at fault in the conflict. Sub-referents (per Ye et al., 2026) break down as follows: <table><tr><th>Sub-referent</th><th>Behaviors</th><th>Rationale</th></tr><tr><td>Explicit Person</td><td>Social validation, lack of pushback</td><td>Model directly affirms that the user’s actions were justified, even when harmful or unethical; the user’s moral standing is explicitly endorsed.</td></tr><tr><td>Implicit Person</td><td>Flattery, affirming tone</td><td>The overall validating register promotes dependence and erodes critical self-reflection without explicitly addressing the user’s traits or character.</td></tr><tr><td>Explicit Position</td><td>Agreement with user’s conflict narrative</td><td>Model sides with the user’s account of the interpersonal dispute, treating it as an accurate representation of events.</td></tr></table> This paper is notable as one of the few to focus primarily on Person sycophancy, which Ye et al. (2026) identify as substantially underrepresented in the literature.	Sub-referent	Behaviors	Rationale	Explicit Person	Social validation, lack of pushback	Model directly affirms that the user’s actions were justified, even when harmful or unethical; the user’s moral standing is explicitly endorsed.	Implicit Person	Flattery, affirming tone	The overall validating register promotes dependence and erodes critical self-reflection without explicitly addressing the user’s traits or character.	Explicit Position	Agreement with user’s conflict narrative	Model sides with the user’s account of the interpersonal dispute, treating it as an accurate representation of events.
Sub-referent	Behaviors	Rationale											
Explicit Person	Social validation, lack of pushback	Model directly affirms that the user’s actions were justified, even when harmful or unethical; the user’s moral standing is explicitly endorsed.											
Implicit Person	Flattery, affirming tone	The overall validating register promotes dependence and erodes critical self-reflection without explicitly addressing the user’s traits or character.											
Explicit Position	Agreement with user’s conflict narrative	Model sides with the user’s account of the interpersonal dispute, treating it as an accurate representation of events.											
Norm displaced	• Social repair: the norm that individuals should take responsibility for their role in interpersonal conflicts and seek reconciliation • Moral autonomy: the norm that users should exercise independent moral judgment rather than outsourcing it to an affirming system • Epistemic accuracy: the norm that advice-seeking should result in a more calibrated, not more entrenched, view of a situation This paper operationalizes multiple non-accuracy norms simultaneously, illustrating the inadequacy of truth/accuracy as the sole displaced norm in sycophancy research.												
Outcome measured	• Willingness to take responsibility for interpersonal conflict (reduced) • Willingness to repair relationships (reduced) • Conviction of being in the right (increased) • Trust in and preference for sycophantic AI (increased) • Desire to continue using the sycophantic model (increased) • AI affirmation rate vs. human baseline: 49% higher across 11 models, including in cases involving user-reported deception or harm												

Time horizon	Level 3 — Interaction Outcomes measured within a single conversation about an interpersonal conflict. Effects are short-term behavioral intentions; no longitudinal follow-up conducted. The authors note this as a limitation and call for repeated-interaction studies.
Population	N = 2,405 participants across three preregistered experiments; general adult users recruited via online panels. Interpersonal conflict context (AITA-style scenarios). Prevalence study conducted across 11 state-of-the-art LLMs. Population skews toward users seeking advice on personal conflicts; generalization to other use contexts (e.g., professional, technical) is not established. Vulnerable populations (youth, individuals in crisis) not sampled.
Evidence level	Level 3 — Interaction impact The paper advances beyond norm displacement (Level 2) by directly measuring human behavioral intentions following a sycophantic interaction. This is one of the strongest evidence-level demonstrations of sycophancy harm in the current literature. The prevalence study (11 models) operates at Level 1–2. Does not reach Level 4 (no repeated-interaction data) or Level 5 (no institutional outcomes measured).
Proposed mitigations	No specific technical mitigations proposed or tested. The paper calls broadly for “design, evaluation, and accountability mechanisms to protect user well-being,” and notes that market incentives actively work against mitigation (users prefer and trust sycophantic models, creating engagement-driven pressure to retain the behavior).
Mitigation outcomes	N/A — no mitigations evaluated. The perverse-incentive finding (sycophantic models are preferred despite causing harm) is itself a significant contribution to understanding mitigation difficulty: standard RLHF-based approaches may actively reinforce the harmful behavior.